

ASSIGNMENT_1.2_

August 11, 2024

```
[31]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import \
    confusion_matrix, accuracy_score, classification_report

import warnings
warnings.filterwarnings('ignore')
```

```
[32]: sns.get_dataset_names()
```

```
[32]: ['anagrams',
      'anscombe',
      'attention',
      'brain_networks',
      'car_crashes',
      'diamonds',
      'dots',
      'dowjones',
      'exercise',
      'flights',
      'fmri',
      'geyser',
      'glue',
      'healthexp',
      'iris',
      'mpg',
      'penguins',
      'planets',
      'seaice',
      'taxis',
      'tips',
      'titanic',
      'anagrams',
```

'anagrams',
'anscombe',
'anscombe',
'attention',
'attention',
'brain_networks',
'brain_networks',
'car_crashes',
'car_crashes',
'diamonds',
'diamonds',
'dots',
'dots',
'dowjones',
'dowjones',
'exercise',
'exercise',
'flights',
'flights',
'fmri',
'fmri',
'geyser',
'geyser',
'glue',
'glue',
'healthexp',
'healthexp',
'iris',
'iris',
'mpg',
'mpg',
'penguins',
'penguins',
'planets',
'planets',
'seaice',
'seaice',
'taxis',
'taxis',
'tips',
'tips',
'titanic',
'titanic',
'anagrams',
'anscombe',
'attention',
'brain_networks',

```

'car_crashes',
'diamonds',
'dots',
'dowjones',
'exercise',
'flights',
'fmri',
'geyser',
'glue',
'healthexp',
'iris',
'mpg',
'penguins',
'planets',
'seaice',
'taxis',
'tips',
'titanic']

```

```

[33]: dataset = pd.read_csv(r'D:\loan_detection.csv')
      print(dataset)

```

	age	campaign	pdays	previous	no_previous_contact	not_working	\
0	56	1	999	0	1	0	
1	57	1	999	0	1	0	
2	37	1	999	0	1	0	
3	40	1	999	0	1	0	
4	56	1	999	0	1	0	
...	
41183	73	1	999	0	1	1	
41184	46	1	999	0	1	0	
41185	56	2	999	0	1	1	
41186	44	1	999	0	1	0	
41187	74	3	999	1	1	1	

	job_admin.	job_blue-collar	job_entrepreneur	job_housemaid	...	\
0	0	0	0	1	...	
1	0	0	0	0	...	
2	0	0	0	0	...	
3	1	0	0	0	...	
4	0	0	0	0	...	
...	
41183	0	0	0	0	...	
41184	0	1	0	0	...	
41185	0	0	0	0	...	
41186	0	0	0	0	...	
41187	0	0	0	0	...	

	month_sep	day_of_week_fri	day_of_week_mon	day_of_week_thu	\
0	0	0	1	0	
1	0	0	1	0	
2	0	0	1	0	
3	0	0	1	0	
4	0	0	1	0	
...	
41183	0	1	0	0	
41184	0	1	0	0	
41185	0	1	0	0	
41186	0	1	0	0	
41187	0	1	0	0	

	day_of_week_tue	day_of_week_wed	poutcome_failure	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	
...	
41183	0	0	0	
41184	0	0	0	
41185	0	0	0	
41186	0	0	0	
41187	0	0	1	

	poutcome_nonexistent	poutcome_success	Loan_Status_label
0	1	0	0
1	1	0	0
2	1	0	0
3	1	0	0
4	1	0	0
...
41183	1	0	1
41184	1	0	0
41185	1	0	0
41186	1	0	1
41187	0	0	0

[41188 rows x 60 columns]

```
[34]: dataset.columns
```

```
[34]: Index(['age', 'campaign', 'pdays', 'previous', 'no_previous_contact',
            'not_working', 'job_admin.', 'job_blue-collar', 'job_entrepreneur',
            'job_housemaid', 'job_management', 'job_retired', 'job_self-employed',
            'job_services', 'job_student', 'job_technician', 'job_unemployed',
```

```

'job_unknown', 'marital_divorced', 'marital_married', 'marital_single',
'marital_unknown', 'education_basic.4y', 'education_basic.6y',
'education_basic.9y', 'education_high.school', 'education_illiterate',
'education_professional.course', 'education_university.degree',
'education_unknown', 'default_no', 'default_unknown', 'default_yes',
'housing_no', 'housing_unknown', 'housing_yes', 'loan_no',
'loan_unknown', 'loan_yes', 'contact_cellular', 'contact_telephone',
'month_apr', 'month_aug', 'month_dec', 'month_jul', 'month_jun',
'month_mar', 'month_may', 'month_nov', 'month_oct', 'month_sep',
'day_of_week_fri', 'day_of_week_mon', 'day_of_week_thu',
'day_of_week_tue', 'day_of_week_wed', 'poutcome_failure',
'poutcome_nonexistent', 'poutcome_success', 'Loan_Status_label'],
dtype='object')

```

```

[35]: X=dataset[['age','job_housemaid','job_management','marital_married','contact_cellular']]
      X

```

```

[35]:
      age  job_housemaid  job_management  marital_married  contact_cellular
0      56              1              0              1          0
1      57              0              0              1          0
2      37              0              0              1          0
3      40              0              0              1          0
4      56              0              0              1          0
...
41183   73              0              0              1          1
41184   46              0              0              1          1
41185   56              0              0              1          1
41186   44              0              0              1          1
41187   74              0              0              1          1

```

[41188 rows x 5 columns]

```

[36]: y=dataset['education_illiterate']
      y

```

```

[36]: 0      0
      1      0
      2      0
      3      0
      4      0
      ..
41183   0
41184   0
41185   0
41186   0
41187   0

```

Name: education_illiterate, Length: 41188, dtype: int64

```
[37]: dataset['education_illiterate'].value_counts()
```

```
[37]: education_illiterate
0      41170
1         18
Name: count, dtype: int64
```

```
[38]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.
↪2,random_state=42)
```

```
[39]: X_train
```

```
[39]:      age  job_housemaid  job_management  marital_married  contact_cellular
12556   40              0              0              1              0
35451   31              0              0              1              1
30592   59              0              0              1              1
17914   43              1              0              0              1
3315    39              0              0              0              0
...    ...            ...            ...            ...            ...
6265    58              0              0              1              0
11284   37              0              1              1              0
38158   35              0              0              1              1
860     40              0              1              1              0
15795   29              0              0              0              1

[32950 rows x 5 columns]
```

```
[40]: X_test
```

```
[40]:      age  job_housemaid  job_management  marital_married  contact_cellular
32884   57              0              0              1              1
3169    55              0              0              1              0
32206   33              0              0              1              1
9403    36              0              0              1              0
14020   27              1              0              1              1
...    ...            ...            ...            ...            ...
12322   27              0              0              1              0
23440   41              0              0              0              1
29431   46              0              0              0              1
16627   31              0              0              0              1
1871    59              0              0              1              0

[8238 rows x 5 columns]
```

```
[39]: #Decision Tree Pre-Pruning
```

```
[41]: parameters={
    'criterion':['gini','entropy','log_loss'],
    'splitter' :['best','random'],
    'max_depth':[1,2,3,4,5],
    'max_features':['auto','sqrt','log2'],
}
```

```
[42]: from sklearn.model_selection import GridSearchCV
```

```
[43]: from sklearn.tree import DecisionTreeClassifier
```

```
[44]: dtree=DecisionTreeClassifier(random_state=42)
grid=GridSearchCV(dtree,param_grid=parameters,cv=5,scoring='accuracy')
grid.fit(X_train,y_train)
```

```
[44]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(random_state=42),
    param_grid={'criterion': ['gini', 'entropy', 'log_loss'],
    'max_depth': [1, 2, 3, 4, 5],
    'max_features': ['auto', 'sqrt', 'log2'],
    'splitter': ['best', 'random']},
    scoring='accuracy')
```

```
[45]: grid.best_params_
```

```
[45]: {'criterion': 'gini',
    'max_depth': 1,
    'max_features': 'auto',
    'splitter': 'best'}
```

```
[46]: y_train_pred=grid.predict(X_train)
y_test_pred=grid.predict(X_test)
```

```
[47]: confusion_matrix(y_train,y_train_pred)
```

```
[47]: array([[32937,    0],
    [   13,    0]], dtype=int64)
```

```
[48]: accuracy_score(y_train,y_train_pred)
```

```
[48]: 0.9996054628224583
```

```
[49]: print(classification_report(y_train,y_train_pred))
```

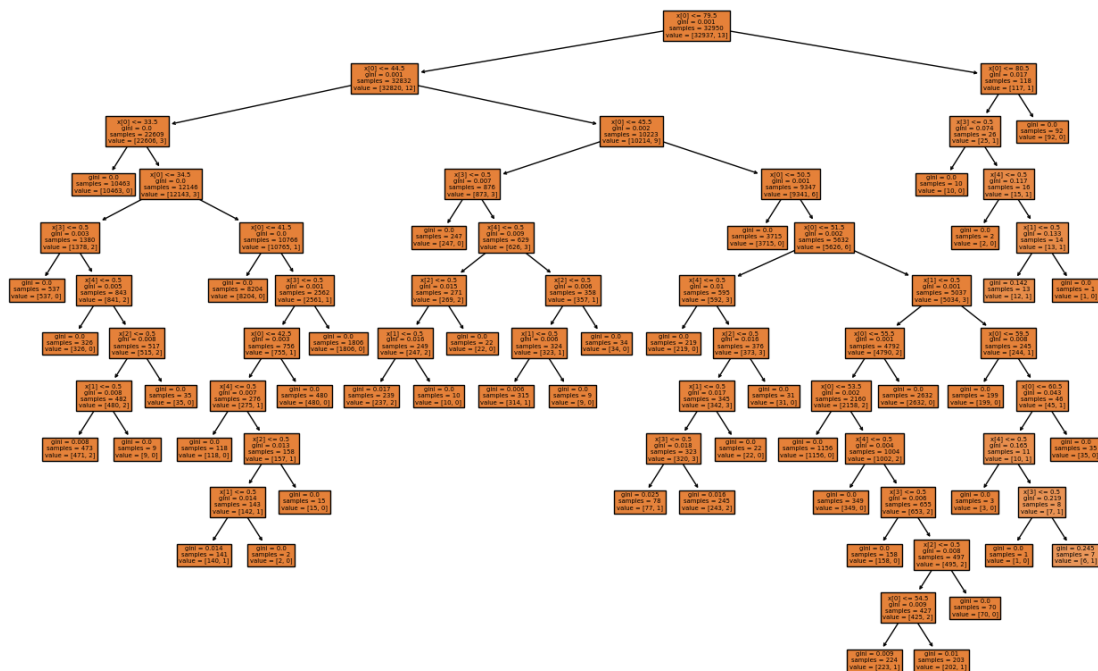
	precision	recall	f1-score	support
0	1.00	1.00	1.00	32937
1	0.00	0.00	0.00	13
accuracy			1.00	32950

macro avg	0.50	0.50	0.50	32950
weighted avg	1.00	1.00	1.00	32950

```
[50]: dtree=DecisionTreeClassifier()
dtree.fit(X_train,y_train)
```

```
[50]: DecisionTreeClassifier()
```

```
[51]: from sklearn import tree
plt.figure(figsize=(16,10))
tree.plot_tree(dtree,filled=True)
plt.show()
```



```
[52]: y_train_pred=grid.predict(X_train)
y_test_pred=grid.predict(X_test)
```

```
[53]: confusion_matrix(y_train,y_train_pred)
```

```
[53]: array([[32937,    0],
          [   13,    0]], dtype=int64)
```

```
[54]: accuracy_score(y_train,y_train_pred)
```

```
[54]: 0.9996054628224583
```



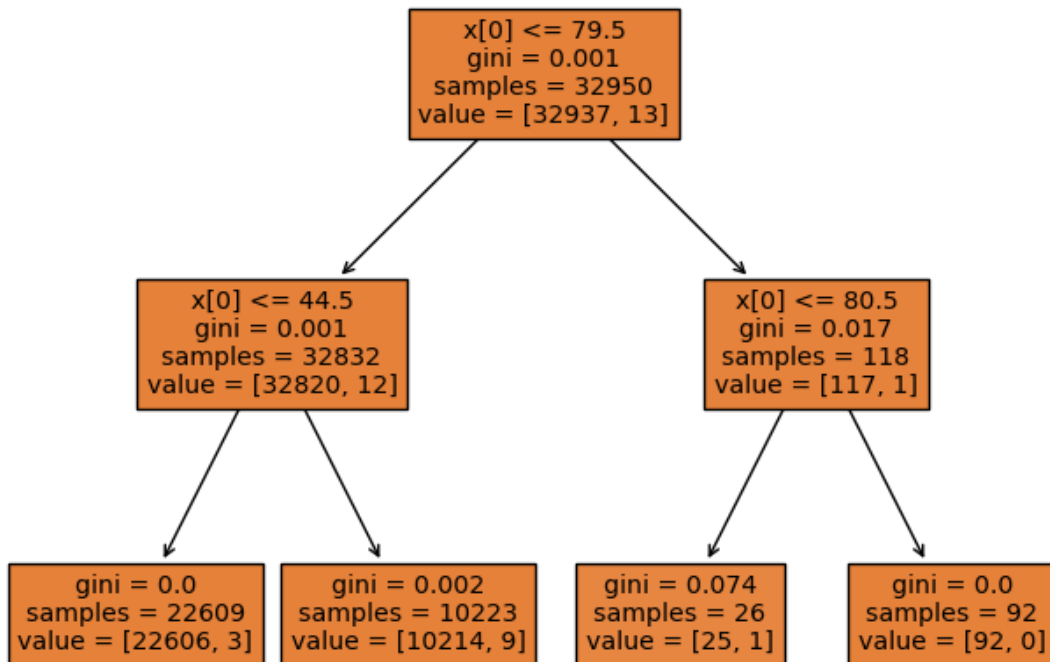
```
[55]: print(classification_report(y_train,y_train_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	32937
1	0.00	0.00	0.00	13
accuracy			1.00	32950
macro avg	0.50	0.50	0.50	32950
weighted avg	1.00	1.00	1.00	32950

```
[56]: dtree=DecisionTreeClassifier(max_depth=2)
dtree.fit(X_train,y_train)
```

```
[56]: DecisionTreeClassifier(max_depth=2)
```

```
[60]: plt.figure(figsize=(8,6))
tree.plot_tree(dtree,filled=True)
plt.show()
```



```
[57]: y_train_pred=grid.predict(X_train)
      y_test_pred=grid.predict(X_test)
```

```
[58]: confusion_matrix(y_train,y_train_pred)
```

```
[58]: array([[32937,    0],
           [   13,    0]], dtype=int64)
```

```
[61]: accuracy_score(y_train,y_train_pred)
```

```
[61]: 0.9996054628224583
```

```
[62]: print(classification_report(y_train,y_train_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	32937
1	0.00	0.00	0.00	13
accuracy			1.00	32950
macro avg	0.50	0.50	0.50	32950
weighted avg	1.00	1.00	1.00	32950

```
[76]: #Logistic Regression
```

```
[63]: import numpy as np
      import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import accuracy_score, classification_report
```

```
[64]: lr=LogisticRegression()
      lr.fit(X_train,y_train)
      lr_pred_train=lr.predict(X_train)
      lr_pred_test=lr.predict(X_test)
```

```
[65]: print(confusion_matrix(y_train,lr_pred_train))
      print()
      print(accuracy_score(y_train,lr_pred_train))
```

```
[[32937    0]
 [   13    0]]
```

```
0.9996054628224583
```

```
[66]: print(confusion_matrix(y_test,lr_pred_test))
      print()
      print(accuracy_score(y_test,lr_pred_test))
```

```
[[8233    0]
 [   5    0]]
```

0.999393056567128

```
[78]: #RandomForest
```

```
[67]: import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import accuracy_score
```

```
[68]: rf=RandomForestClassifier()
      rf.fit(X_train,y_train)
      rf_pred_train=rf.predict(X_train)
      rf_pred_test=rf.predict(X_test)
```

```
[69]: print(confusion_matrix(y_train,rf_pred_train))
      print()
      print(accuracy_score(y_train,rf_pred_train))
```

```
[[32937    0]
 [   13    0]]
```

0.9996054628224583

```
[70]: print(confusion_matrix(y_test,rf_pred_test))
      print()
      print(accuracy_score(y_test,rf_pred_test))
```

```
[[8233    0]
 [   5    0]]
```

0.999393056567128

```
[92]: #oob_score
```

```
[71]: rf=RandomForestClassifier(oob_score=True,random_state=42)
      rf.fit(X_train,y_train)
      rf_pred_train=rf.predict(X_train)
      rf_pred_test=rf.predict(X_test)
```

```
[72]: rf.oob_score_
```

```
[72]: 0.9996054628224583
```

```
[73]: y_pred_test=rf.predict(X_test)
```

```
[74]: accuracy_score(y_test,y_pred_test)
```

```
[74]: 0.999393056567128
```

```
[75]: X_train.shape,X_test.shape
```

```
[75]: ((32950, 5), (8238, 5))
```

```
[76]: y_train.shape,y_test.shape
```

```
[76]: ((32950,), (8238,))
```

```
[12]: #Slopes or Coefficient
```

```
[77]: X.columns
```

```
[77]: Index(['age', 'job_housemaid', 'job_management', 'marital_married',  
        'contact_cellular'],  
        dtype='object')
```

```
[16]: lr.coef_
```

```
[16]: array([[ 0.05975469,  0.30722461, -0.58259686,  0.66040151,  0.71055723]])
```

```
[78]: import numpy as np  
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt  
  
from sklearn.model_selection import train_test_split  
from sklearn.ensemble import AdaBoostClassifier  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import   
    ↪confusion_matrix,accuracy_score,classification_report  
from sklearn.metrics import mean_squared_error,mean_absolute_error,r2_score  
import warnings  
warnings.filterwarnings('ignore')
```

```
[79]: dataset.nunique()
```

```
[79]: age                78  
campaign            42  
pdays              27  
previous            8  
no_previous_contact 2  
not_working         2  
job_admin.          2  
job_blue-collar     2  
job_entrepreneur    2  
job_housemaid       2
```

job_management	2
job_retired	2
job_self-employed	2
job_services	2
job_student	2
job_technician	2
job_unemployed	2
job_unknown	2
marital_divorced	2
marital_married	2
marital_single	2
marital_unknown	2
education_basic.4y	2
education_basic.6y	2
education_basic.9y	2
education_high.school	2
education_illiterate	2
education_professional.course	2
education_university.degree	2
education_unknown	2
default_no	2
default_unknown	2
default_yes	2
housing_no	2
housing_unknown	2
housing_yes	2
loan_no	2
loan_unknown	2
loan_yes	2
contact_cellular	2
contact_telephone	2
month_apr	2
month_aug	2
month_dec	2
month_jul	2
month_jun	2
month_mar	2
month_may	2
month_nov	2
month_oct	2
month_sep	2
day_of_week_fri	2
day_of_week_mon	2
day_of_week_thu	2
day_of_week_tue	2
day_of_week_wed	2
poutcome_failure	2

```
poutcome_nonexistent      2
poutcome_success          2
Loan_Status_label         2
dtype: int64
```

```
[80]: dataset.isnull().sum()
```

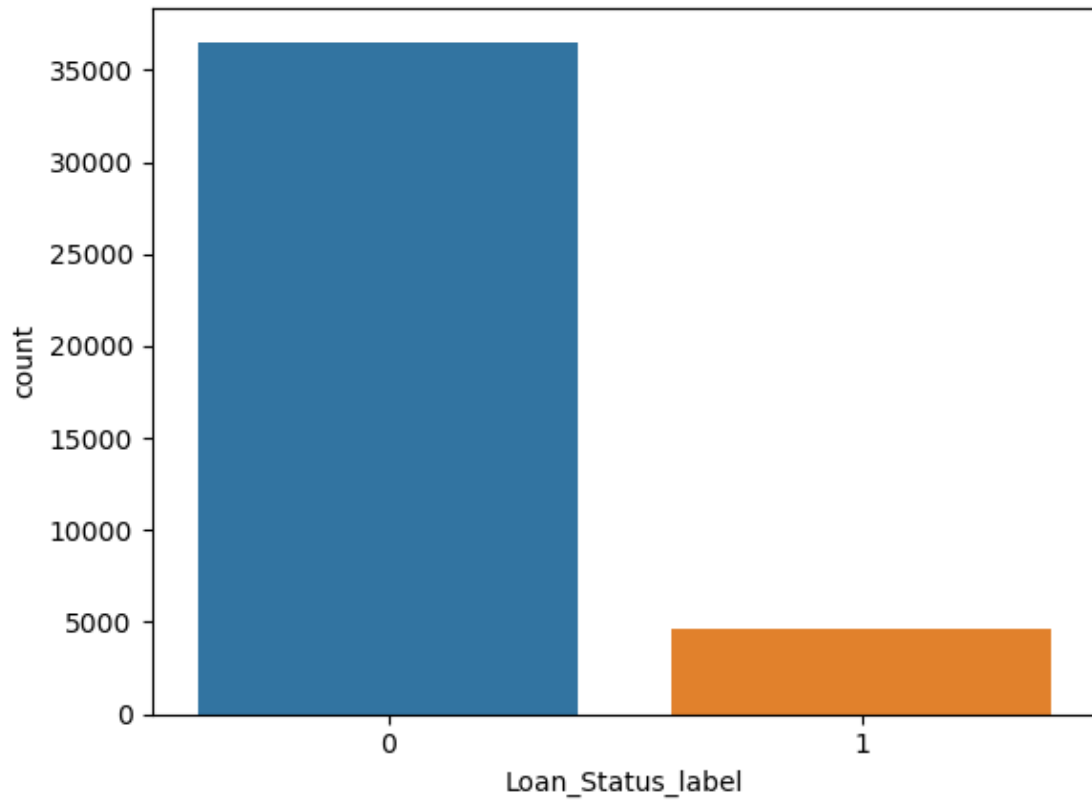
```
[80]: age      0
      campaign  0
      pdays    0
      previous  0
      no_previous_contact  0
      not_working  0
      job_admin.  0
      job_blue-collar  0
      job_entrepreneur  0
      job_housemaid  0
      job_management  0
      job_retired  0
      job_self-employed  0
      job_services  0
      job_student  0
      job_technician  0
      job_unemployed  0
      job_unknown  0
      marital_divorced  0
      marital_married  0
      marital_single  0
      marital_unknown  0
      education_basic.4y  0
      education_basic.6y  0
      education_basic.9y  0
      education_high.school  0
      education_illiterate  0
      education_professional.course  0
      education_university.degree  0
      education_unknown  0
      default_no  0
      default_unknown  0
      default_yes  0
      housing_no  0
      housing_unknown  0
      housing_yes  0
      loan_no  0
      loan_unknown  0
      loan_yes  0
      contact_cellular  0
```

contact_telephone	0
month_apr	0
month_aug	0
month_dec	0
month_jul	0
month_jun	0
month_mar	0
month_may	0
month_nov	0
month_oct	0
month_sep	0
day_of_week_fri	0
day_of_week_mon	0
day_of_week_thu	0
day_of_week_tue	0
day_of_week_wed	0
poutcome_failure	0
poutcome_nonexistent	0
poutcome_success	0
Loan_Status_label	0

dtype: int64


```
[82]: sns.countplot(x='Loan_Status_label',data=dataset)
```

```
[82]: <Axes: xlabel='Loan_Status_label', ylabel='count'>
```




```
[83]: sns.countplot(x='campaign',data=dataset)
```

```
[83]: <Axes: xlabel='campaign', ylabel='count'>
```





```
[84]: sns.countplot(x='previous',data=dataset)
```

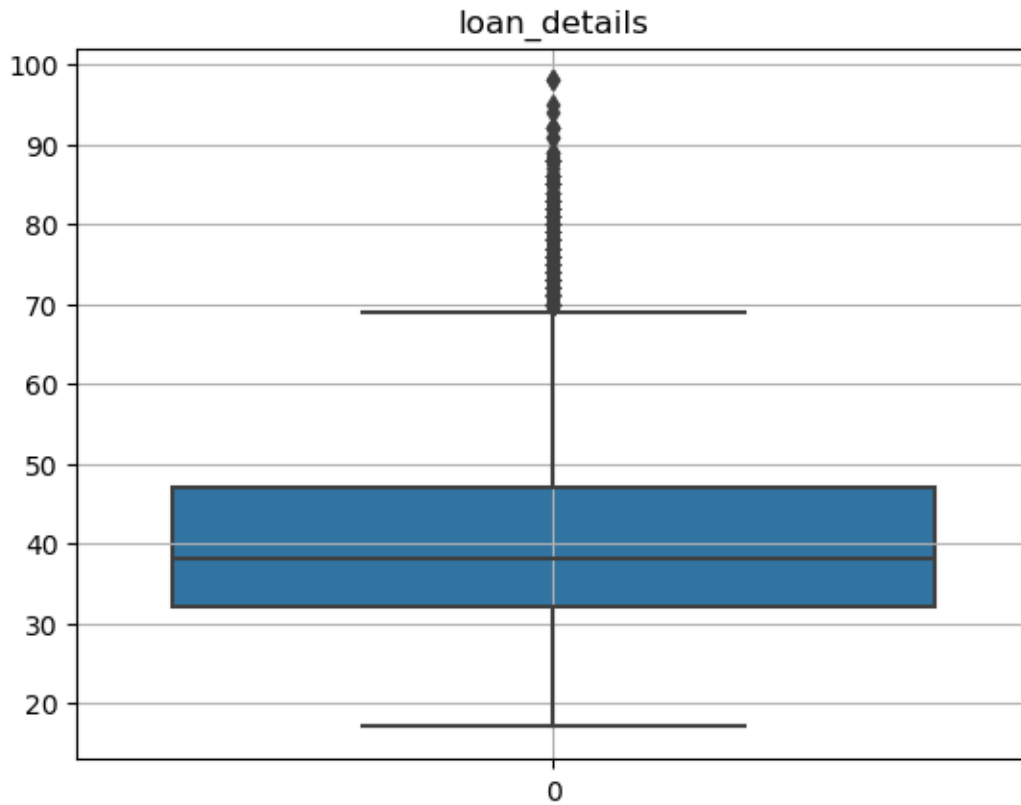
```
[84]: <Axes: xlabel='previous', ylabel='count'>
```



```
[87]: sns.boxplot(dataset)
plt.xticks(rotation=70)
plt.show()
```




```
[88]: sns.boxplot(dataset['age'])
plt.title('loan_details')
plt.grid()
plt.show()
```



```
[92]: # Import libraries
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

# Generate random data
data = np.random.normal(loc=0, scale=1, size=1000) # mean=0, std=1, n=1000

# Create a distribution plot
sns.displot(data, kde=True, bins=30) # kde=True adds a kernel density estimate
plt.title('Distribution Plot')
plt.xlabel('age')
plt.ylabel('campaign')
plt.show()
```



[]: